

Basic Architectures

The building blocks for complex systems


Contents

- *Monolith* – a cohesive codebase.
 - *Shards* – multiple instances of a (sub)system.
 - *Layers* – subdivision by the level of abstractness.
 - *Services* – components, dedicated to subdomains.
 - *Pipeline* – a chain of data processing steps.
- ... and common variants of each of the architectures.

A bit of background

We are comparing patterns by drawing them with common axes:

- Abstractness
- Subdomain
- Sharding




A basic architecture is partitioned along one of these axes

Monolith

A cohesive (sub-)system

Monolith – overview



*Monolith,
Big Ball of Mud*

A *Monolith* (“single stone”) is a cohesive (sub-)system.

- In a strict sense, it cannot be subdivided.
- In a loose sense, we don’t want to look inside it.

The word “*Monolith*” has been redefined several times (the latest version being “a single unit of deployment”) but we’ll stick to its original meaning.

Monolith is good for quickly coding a tiny one-off project.

Anything larger or longer-living will be difficult to support.

Monolith – trade-offs

Is easy to write and debug

Allows for low latency

Permits thorough performance optimizations

Maintains a self-consistent state

Is hard to read, change, and support

Does not scale

Is not fault tolerant

Is developed by a single team

Is limited to one technology and style


Aligns with a single set of forces

Monolith – misnomers


A few non-monolithic architectures
are often called “monoliths”:

Name	Meaning
Monolithic Lambda (<i>Lambdalith</i>)	Stateless <i>Shards</i>
Layered Monolith	<i>Layers</i>
Modular Monolith (<i>Modulith</i>)	Co-located <i>Services</i>
Distributed Monolith	Usually <i>Service-Oriented Architecture</i>


Monolith – event handling



A blocking thread per request




A single non-blocking thread processes all events



A coroutine or fiber per request

Monolith – event handling




A *Reactor* runs each task in a blocking thread.

- A backend *Reactor* uses OS threads to execute many requests in parallel. Any shared resources must be protected with locks.
- An embedded system may wrap a physical device with a single-threaded *Reactor* to avoid race conditions in the hardware.

Reactor's code is simple & stupid.

However, each thread takes a lot of system resources.

Monolith – event handling




A *Proactor* handles everything with a single non-blocking thread.

There is no need for locks and no associated delays, therefore the system reacts to events in real time and its state is always up-to-date.

Proactor fits real-time control systems.

Sadly, its logic is dispersed over many event handlers.

Monolith – event handling



*Half-Sync /
Half-Async*

A *Half-Sync / Half-Async* system allocates a blocking coroutine for each task. All the coroutines run in the same thread or thread pool.

- The upper half of the system is like a *Reactor*.
- The lower half is a *Proactor*.


The code of each coroutine is simple.

But the underlying framework is complex.

Shards

Multiple instances of a (sub-)system

Shards – overview



*Shards, Instances,
Replicas*

Shards are multiple instances of a component.

They may or may not differ in their data.

In most systems shards don't intercommunicate directly.

Shards provide scalability and fault tolerance...

...as long as there is no shared data.

Shards – trade-offs

The (sub-)system becomes scalable

Sharding may be used to improve latency and / or fault tolerance


It is hard to achieve state consistency

There is an operational effort to deploy or update the shards

Shards – isolation

Isolation	Benefits	Drawbacks
A thread per shard	Limited scalability	Shared data needs locks
A process per shard	Software fault tolerance	Sharing data is non-trivial
A server per shard	Full scalability and fault tolerance	Sharing data is impossible State synchronization is complicated

Shards – state



Shards, Partitions

Each *partition* owns a slice of the system's data.


A client's access is usually limited to one *partition*.

A *Sharding Proxy* connects clients to their *partitions*.

Partitioning scales the data storage.

No data can be shared among the system's clients.

Shards – state



Each *replica* has a copy of the whole system's data.

A *Load Balancer* connects new clients to idle *replicas*.


The *replicas* need to keep their data in sync:

- There may be a *leader* which processes all write requests and broadcasts the changes to the *followers*.
- Or the data to be changed must be locked in all the *replicas* to avoid write conflicts.

Replication improves fault tolerance and read performance.

But it is costly and complicates write access to the data.

Shards – state



System *instances* may be stateless.

That requires both a *Load Balancer* to choose an *instance* from the *pool* and a *Shared Repository* to persist the data.

The *Shared Repository* (database or file storage) may become the single point of failure and performance bottleneck.


Instances, Lambdas

The number of instances in the *pool* may be fixed or elastic.

Stateless instances are perfectly scalable and fault tolerant.

But now the *Shared Repository* becomes the bottleneck.

Shards – state



A stateful instance of a component can be created for each client.

Such an *Actor* keeps the client's data readily available in its memory and acts on behalf of its client.

The changes to *Actors'* states are persisted to a database.


Actors grant low request latency and fault isolation between clients.

Complex logic that involves multiple clients is impossible.

Layers

Components that differ in abstractness

Layers – overview



Layers are system-wide components that differ in their abstractness.

An upper *layer* may depend on a lower *layer* or its interface. A lower *layer* knows nothing about the *layers* above it.

Tiers are distributed *Layers*.

Layering provides flexibility to small- and medium-sized projects...

...at the cost of lost optimization opportunities.

Layers – trade-offs

Each layer may be developed by a dedicated team


The layers may differ in technologies, scalability, security, and even hardware and physical location

The business logic is separated from reusable generic code

The number of layers and associated teams is usually limited to 3 or 4.


The interfaces between layers tend to negatively impact performance

Layers – openness



An *open layer* is transparent – the layer above it may directly access the layer below it.


A *closed layer* is opaque – it hides everything below it from its clients.



Layers – isolation

Isolation	Benefits	Drawbacks
Synchronous layers	Encapsulation Multiple teams	Lost optimization opportunities
Asynchronous layers	Low latency	Complicated debugging
A process per layer	Independent technologies, deployment and scalability Software fault isolation	Inconsistent state after errors
Distributed tiers	Independent hardware and location	Slow communication between layers

Layers – examples




DDD Layers

Domain-Driven Design defines the following layers:

- *Presentation* (user interface)
- *Application* (integration or use cases)
- *Domain* (business rules)
- *Infrastructure* (utilities and data access)

Layers – examples



3-Tier system

Tiers are distributed Layers.

3-Tier Architecture contains:


- *Frontend* (user interface)
- *Backend* (business logic)
- *Database* (persistent data)

The tiers differ in their scalability, location, and security.

Services

Subdomain-aligned components

Services – overview



Services, Modules

Services employ a component per subdomain.

When the subdomains are loosely coupled, the corresponding services can be developed by fairly independent teams with minimal communication overhead.

Services are good for large projects with several teams.

They require the subdomains to be stable and loosely coupled.

Services – trade-offs

Services fit large projects with several teams

The subdomains may vary in technologies and scaling

A degree of fault tolerance is achieved with loosely coupled services

Use cases that involve multiple subdomains become slow and complicated

It is hard to synchronize the states of the services

Subdomains boundaries should never change

There is operational complexity


Services – isolation

Isolation	Benefits	Drawbacks
Synchronous modules	Multi-team development	Subdomain boundaries are frozen
Asynchronous modules	Event replay	Hard to share data or debug across services
Multiple processes	Independent technologies and scalability	Inconsistent state after errors
Distributed services	Granular scalability Good fault isolation	High communication overhead


Services – size

Size	Architecture	Traits
Whole subdomain	Service-Based Architecture	Easy to design and implement
Part of a subdomain	Microservices	Fine-grained scalability
Class-like	Actors	Real-time latency and low resource consumption
Single function	Nanoservices	Reusable components?

Services – internals



Services – examples




*Service-Based
Architecture*

Service-Based Architecture (SBA) is made of coarse-grained subdomain services.

Some or all of them may share a database. They may not be independently deployable.

Being pragmatic, *SBA* violates best practices in favor of simplicity.

Services – examples




Microservices

Microservices follow best practices:

- Private code and databases.
- Independent scaling and deployment with a *Service Mesh*.
- Mostly asynchronous communication.
- Fine-grained responsibilities.

Shared libraries are placed into *Sidecars*.

Services – examples



Actors

Actors are asynchronous objects.


Each actor represents a physical or logical entity, such as a bank client or chat participant.

An *Actor Framework* runs millions of interconnected actors.

Pipeline

A component per step of data processing

Pipeline – overview



Pipeline

A *Pipeline* passes a stream of data through a chain of components that transform it or react to it.

As each component has well-defined input(s) and output(s), it is easy to test in isolation or replace.

Pipeline is a kind of *Services* with unidirectional data flow.

Pipeline fits simple data-intensive projects.

Any integration logic multiplies system components.

Pipeline – trade-offs

Pipeline supports multiple development teams and technologies

The components are easy to add, remove, replace, and can be tested in isolation


The system is highly scalable

Pipeline supports only a few simple use cases with rudimentary error handling

The latency is high

Not every domain can be represented as a Pipeline

Pipeline – examples




Pipes and Filters

Pipes and Filters is the simplest kind of *Pipeline*. It is often run locally.

A *filter* receives an input, transforms it, and produces an output.

A *pipe* connects the output of a filter to the input of another filter.

Pipeline – examples




*Choreographed
Event-Driven Architecture*

Event-Driven Architecture (EDA) is a tree of subdomain services which subscribe to each other's events.

It is easy to extend with new services.

Pipeline – examples




Data Mesh

Data Mesh is a graph of streams of analytical information. Its nodes are called *Data Product Quanta (DPQ)*.

It extends any kind of distributed system, extracting its data analytics aspect into an overlapping set of services and databases.

Conclusion


Building complex architectures



Orchestrated
Services

Shards, Layers and Services can be combined, often recursively, for example:

- A single *integration layer* over *services* makes *Orchestrated Services*.
- *Layers*, divided into *Services*, make *Service-Oriented Architecture*.



Service-Oriented
Architecture

Links

All of this and much more can be found in my book *Architectural Metapatterns*:

- read online at metapatterns.io
- download from leanpub.com/metapatterns

The diagrams and the ODT source file are available under the CC BY license.

The book is free.
No strings attached.